



# Ethereum Smart Contract Upgradability Workshop

TIEC - Cairo  
April 14th, 2019



Ahmed Ali <[aabdulwahed](#)>



ocean

- Motivation
- Limitations
- Contract Structure
- EVM Data Locations
- Low Level Calls
- Upgradability Patterns
- Recommendations
- Tools: ZeppelinOS
- Discussion/networking

- Smart contracts are immutable code
- Any code is prone to errors and vulnerabilities
- Upgradability means the ability to upgrade smart contracts after they have been deployed!


- Any old deployed contract will stay as is forever on Ethereum
- Users might not know about the release of a new contract version

- **State variables:**

- Permanently stored variables in contract storage.
- A special type of state variables is *constant* state variable.

- **Events:**

- EVM logging facility
- Stored in the transaction log
- Search optimization (indexed)
- Support anonymous events




- **Functions:**

- Executable units of code
- Visibility (External, Public, internal, private)
- constructor a special function, executed only once upon the contract creation

- **Modifiers:**

- Modify the function behavior by automatically check a condition prior to function execution



# Ethereum Smart Contract Upgradability

## Contract Structure

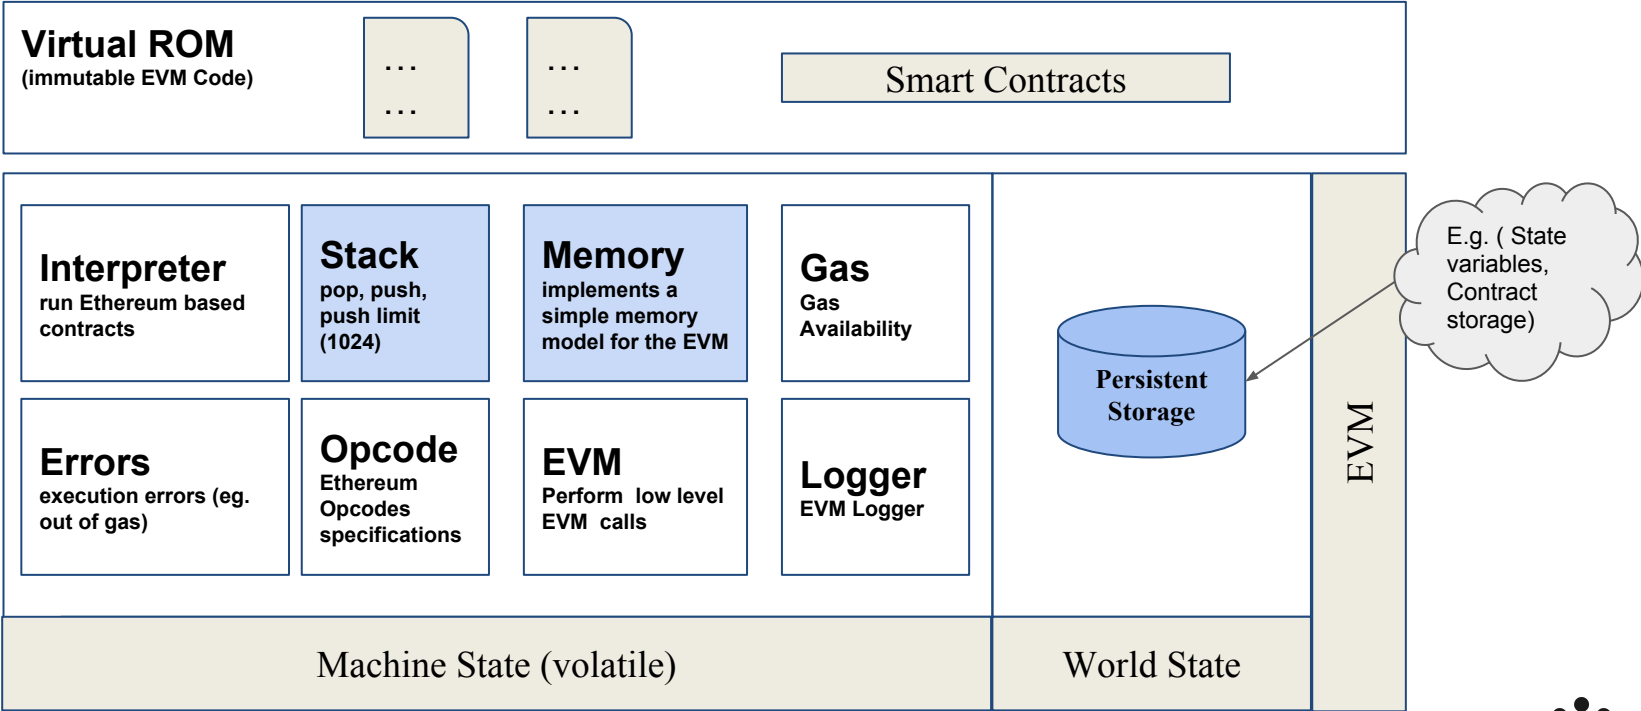
```
contract SampleContract is Ownable {  
    // state variables  
    uint256 public value = 0;  
    // modifier  
    modifier onlyValidValue(uint256 _value)  
    {  
        require(  
            value > 0,  
            'Invalid value'  
        );  
        _;  
    }  
    // event  
    event ValueUpdated(  
        uint256 indexed oldValue,  
        uint256 indexed newValue,  
        uint256 updatedAt  
    );  
};
```

setValue() will  
never work and  
always revert with  
'Invalid value'

```
// constructor  
constructor(address contractOwner)  
    public  
{  
    require(contractOwner != address(0));  
    transferOwnership(contractOwner);  
}  
  
function setValue(uint256 _value)  
    public // visibility is public  
    onlyOwner() // inherited modifier from Ownable!  
    returns(bool)  
{  
    return _setValue(_value);  
}  
  
function _setValue(uint256 newValue)  
    private  
    onlyValidValue(newValue)  
    returns(bool)  
{  
    uint256 oldValue = value;  
    value = newValue;  
    emit ValueUpdated(oldValue, newValue, block.number);  
}
```

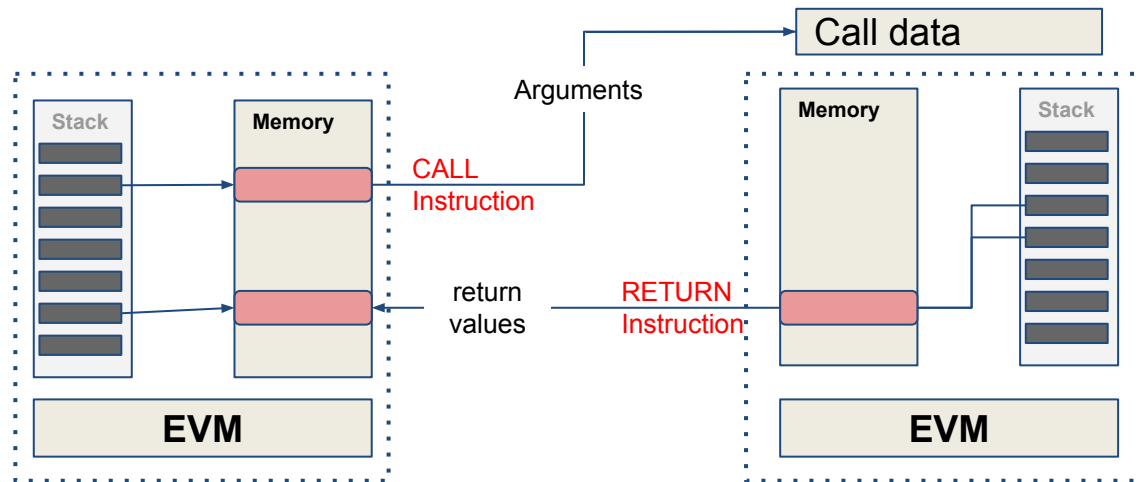
[https://github.com/aabdulwahed/contract-upgradability-Cairo-workshop-2019/tree/master/labs/sample\\_contract](https://github.com/aabdulwahed/contract-upgradability-Cairo-workshop-2019/tree/master/labs/sample_contract)





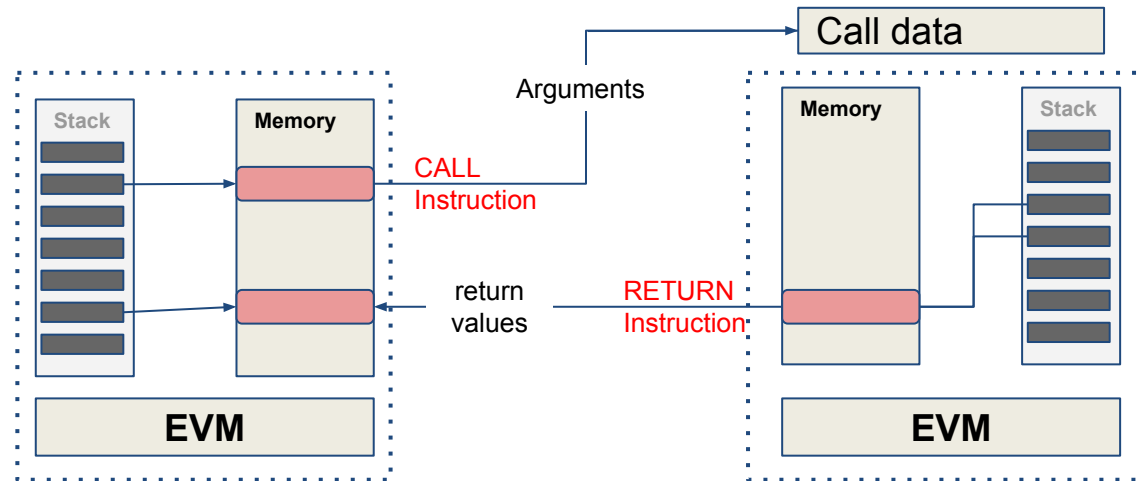


- **Storage:**
  - Persistent read-write word-addressable space
  - Storage is a key-value store that maps 256-bit words to 256-bit words.
  - Access with *SSTORE*/*SLOAD* instructions
  - All locations in storage are well-defined initially as zero
  - *SLOAD* loads a word from storage into the stack
  - *SSTORE* saves a word to storage
- **How does storage allocation work?**
  - Statically sized variables
    - They are laid out contiguously in storage starting from 0
  - Mapping and Dynamically-sized Arrays
    - Starts at unfilled slot in the storage at some position  $P$
    - For arrays, the data is stored in  $\text{keccak256}(p)$
    - For mappings, for key at position  $P$ , the value is stored in  $\text{keccak256}(k \cdot p)$



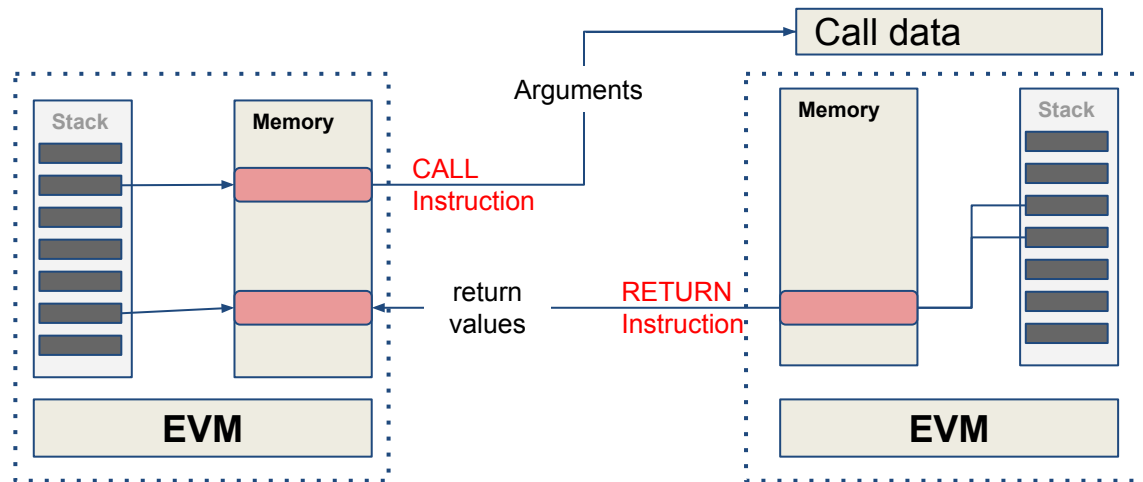
- **Stack:**

- EVM is a 256-bit word machine
- The stack has a maximum size of 1024
- All EVM operations are performed on the stack
- Operations are represented by opcodes (POP, PUSH, DUP, SWAP, ect.)
- The depth of transaction invocation (message call) is limited to less than 1024 levels



- **Memory:**

- Read-write byte-addressable space (word size: 256 bits)
- Memory is linear and can be addressed at byte level
- Access with MSTORE/MSTORE8/MLOAD instructions
- All locations in memory are predefined as zero



- **Calldata:**

- Read-only byte-addressable space
- Holds data parameters of a transaction or call
- Unlike stack, in order to read this data, you have to specify byte offset and number of bytes

- Play with Data locations (inline assembly example)

```
contract PlayWithInLineAssembly {
function add(uint256 _a, uint256 _b) public pure
returns (uint256 result)
{
    assembly {
        // Solidity always stores a free memory pointer at position 0x40
        // load into stack from memory @0x40
        let aPtr := mload(0x40)
        // increment bPtr by adding 32 bytes offset to 0x40
        let bPtr := add(aPtr, 32)
        // copy call data (_a) into memory: after first 4 bytes (function selector)
        calldatacopy(aPtr, 4, 32)
        // copy call data (_b) into memory: after first (4bytes + 32 bytes)
        calldatacopy(bPtr, add(4, 32), 32)
        // load data (aPtr, bPtr) from memory into stack
        result := add(mload(aPtr), mload(bPtr)) // sum two data values and assign the output to result
    }
}
```

Special Low level functions in solidity: *call* vs *Delegatecall*

```
contract calleeContract{
    event ContractAddress(address _from);
    function triggerCall() public payable {
        emit ContractAddress(this);
    }
}

contract delegateCallvsCall {
    function myDelegateCall(address _calleeContract) public returns(bool){
        // emits the address of the callee contract
        require(_calleeContract.call(bytes4(keccak256('triggerCall()'))),
            'invalid low level call');
        // emit the address of this contract address (caller)
        require(_calleeContract.delegatecall(bytes4(keccak256('triggerCall()'))),
            'invalid low level delegatecall');
        return true;
    }
}
```

Special Low level functions in solidity: *call* vs *Delegatecall*

```
contract calleeContract{
```

The screenshot shows a Solidity IDE with the following code:

```
contract calleeContract{  
    ▶ calleeContract at 0x35e...450cf (memory)    
    ▼ delegateCallvsCall at 0xbde...4ddc9 (memory)    
    myDelegateCall "0x35ef07393b57464e93deb59175ff72e6499450cf" 
```

```
        // emits the address  
        require(_calleeContract == this, "not the same");  
        // emit the address  
        require(_calleeContract == this, "not the same");  
        return true;  
    }  
}
```

```
"event": "ContractAddress",  
"args": {  
    "0": "0x35eF07393b57464e93dEB59175fF72E6499450cF",  
    "_from": "0x35eF07393b57464e93dEB59175fF72E6499450cF",  
    "length": 1  
}  
  
"from": "0xbde95422681e4c3984635af2f2f35f8c44a4ddc9",  
"topic": "0x1a0f921ce3c6f2f0f6be5b624a487bcd5143e1fd1833154f39ab63e13d897"  
  
"event": "ContractAddress",  
"args": {  
    "0": "0xBde95422681e4C3984635Af2f2F35f8c44A4ddc9",  
    "_from": "0xBde95422681e4C3984635Af2f2F35f8c44A4ddc9",  
    "length": 1  
}
```

### Two families of patterns

#### Data Separation

This family relies on separation of logic and data. The logic contract is the only authorized contract that can call the data contract


#### Delegatecall-based proxies

Data and logic contracts are kept separate but the proxy contract (data contract) calls the logic contract via *delegatecall*



### Data Separation

- The design is simple, does not require any low level expertise
- Only the owner can alter its content
- For upgradability, we need to understand **how to store data**, and **how to perform the upgrade**.



### Data Separation

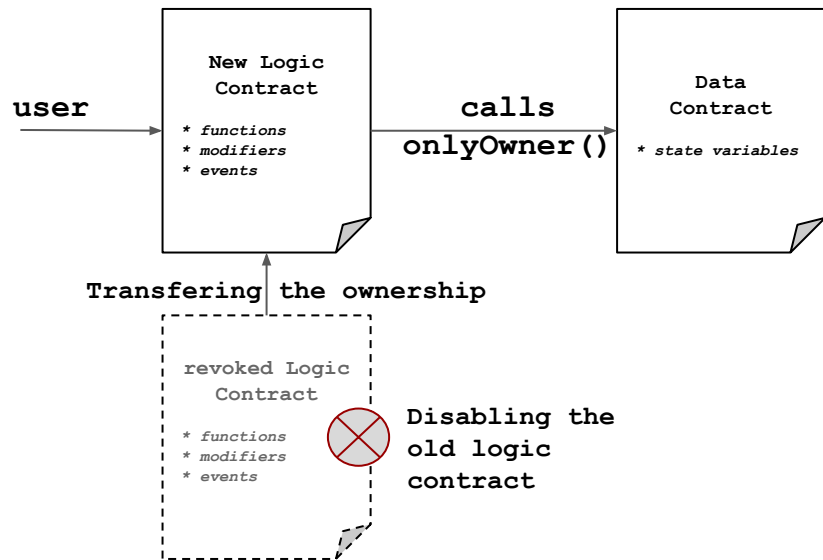
- **How to store data?**
  - Eternal Storage pattern
  - Unified key-value data storage pattern
  - A mapping from a *bytes32* key value to each base variable type

Mappings	
Key	Value
bytes32	uint256
bytes32	int8
bytes32	string

```
contract EternalStorageDataContract is Ownable {  
  
    mapping(bytes32 => uint256) uInt256Storage;  
  
    constructor(address contractOwner)  
    public  
    {  
        require(contractOwner != address(0));  
        transferOwnership(contractOwner);  
    }  
  
    function getUint256(bytes32 key)  
    public  
    view  
    returns(uint256)  
    {  
        return uInt256Storage[key];  
    }  
  
    function setUint256(bytes32 key, uint new_val)  
    public  
    onlyOwner  
    {  
        uInt256Storage[key] = new_val;  
    }  
}
```

### Data Separation

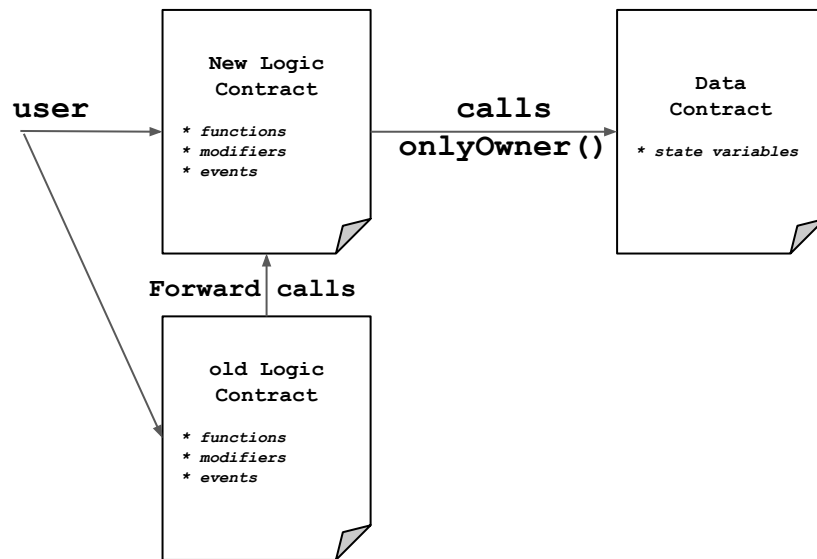
- How to perform upgrade?
  - Mechanisms
    - Puasable mechanism by deploying a new logic contract and transferring the ownership of the data contract to it



**Pausable Mechanism**

### Data Separation

- **How to perform upgrade?**
  - Mechanisms
    - Different approach by forwarding the calls from the original logic contract to the new version




**Forwarding Mechanism**

### Data Separation

- **How to perform upgrade?**

- Mechanisms

- More complex approach, upgrade by deploying a proxy contract that calls a new logic contract.



It provides distinction of responsibilities but it comes with extra gas cost

**Proxy Mechanism**


### Data Separation

- **Risks**

- Adds more complex authorization schema to the code
- Eternal storage increases the complexity of the data model.
- Some developers might implement this pattern incorrectly (e.g keeping some logic in data contract which is impossible to upgrade)

### Delegatecall-based proxies

- Similarly to Data Separation approach
  - This approach splits the contract into two contracts:
    - Proxy contract which holds data
    - Logic contract holds logic
  - But the proxy contract (data contract) calls the logic contract using ***delegatecall*** in the context of the proxy contract



### Delegatecall-based proxies

Bad delegatecall usage  
(Expect memory corruption)

```
contract SampleLogicContract {
    uint public a;

    function set(uint val)
        public
        returns (bool)
    {
        a = val;
        return true;
    }
}
```

```
contract ProxyContract {
    address public contractPtr;
    // different state variable with the same name
    uint public a;


    constructor()
        public
    {
        contractPtr = address(new SampleLogicContract());
    }



    function set(uint val)
        public
        returns (bool)
    {
        bool state = contractPtr.delegatecall(bytes4(keccak256("set(uint256)")), val);
        require(
            state,
            'invalid delegatecall'
        );
        return true;
    }
}
```



### Delegatecall-based proxies

Bad delegatecall usage  
(Expect memory corruption)

Deployed Contracts 

▼ ProxyContract at 0x692...77b3a (memory)  


set 1 ▼



a

0: uint256: 0

contractPtr

0: address: 0x0001

Deployed Contracts 

▼ ProxyContract at 0x692...77b3a (memory)  

set 2 ▼

a

0: uint256: 0

contractPtr

0: address: 0x0001

- SampleLogicContract.set is executed within the context of ProxyContract
- SampleLogicContract knows only one state variable (a)
- If we try to execute set function in the context of ProxyContract, the delegate call will write the value to the first state variable `contractPtr` instead of `a`

### Delegatecall-based proxies

```
contract SampleLogicContract {  
    uint public a;  
  
    function set(uint val)  
        public  
        returns (bool)  
    {  
        a = val;  
        return true;  
    }  
}
```

### Safe delegatecall proxy

```
contract ProxyContract {  
    uint public a;  
    address public contractPtr;  
  
    constructor()  
        public  
    {  
        contractPtr = address(new SampleLogicContract());  
    }  
  
    function set(uint val)  
        public  
        returns (bool)  
    {  
        bool state = contractPtr.delegatecall(bytes4(keccak256("set(uint256)")), val);  
        require(  
            state,  
            'invalid delegatecall'  
        );  
        return true;  
    }  
}
```

Delegatecall-based proxies

Safe delegatecall proxy

ProxyContract at 0x089...659fb (memory)

set

1

▼

a

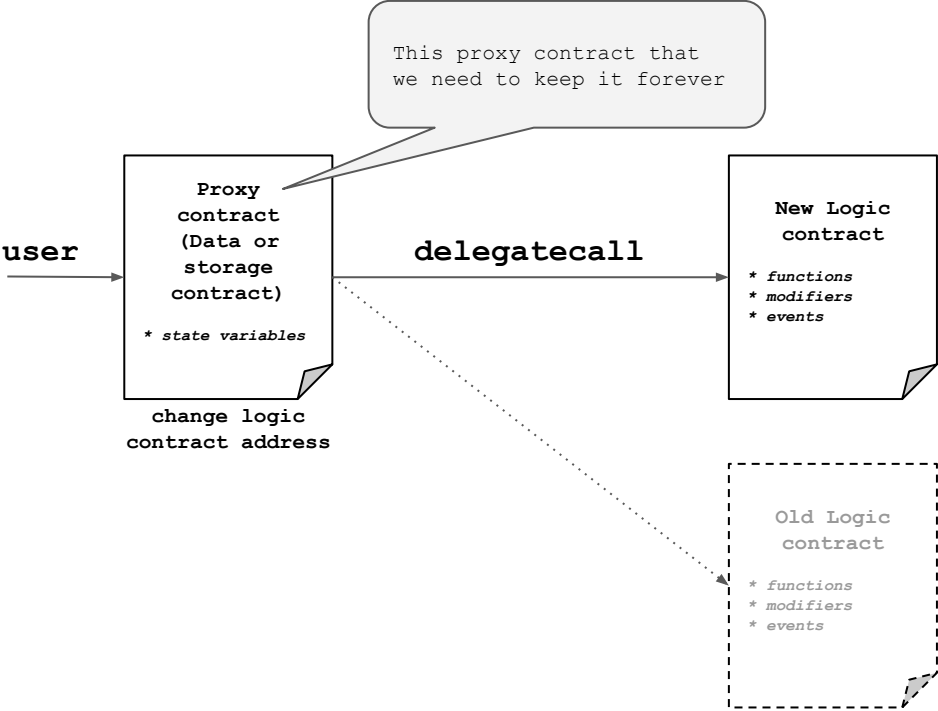
0: uint256: 1

contractPtr

0: address: 0x0dBBc8566E6aaA8bE81dE21850cd177F88b8d648

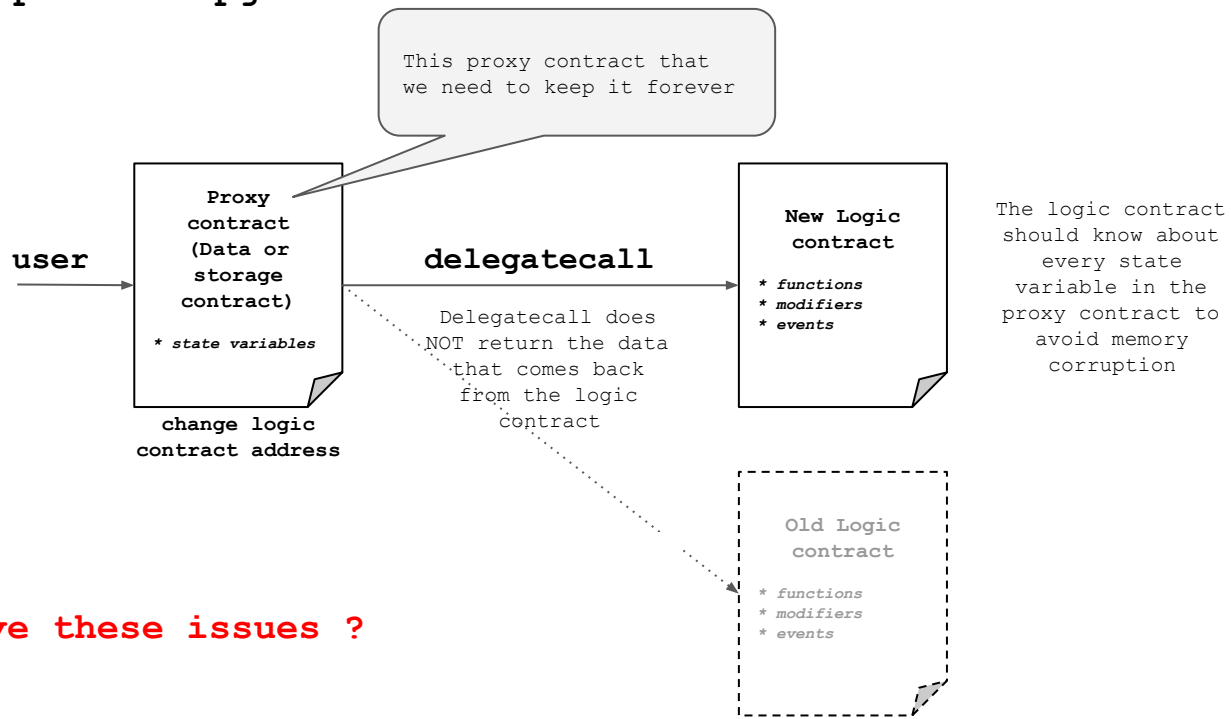
### Delegatecall-based proxies

- How to perform upgrade



### Delegatecall-based proxies


- How to perform upgrade



How to solve these issues ?

### Delegatecall-based proxies

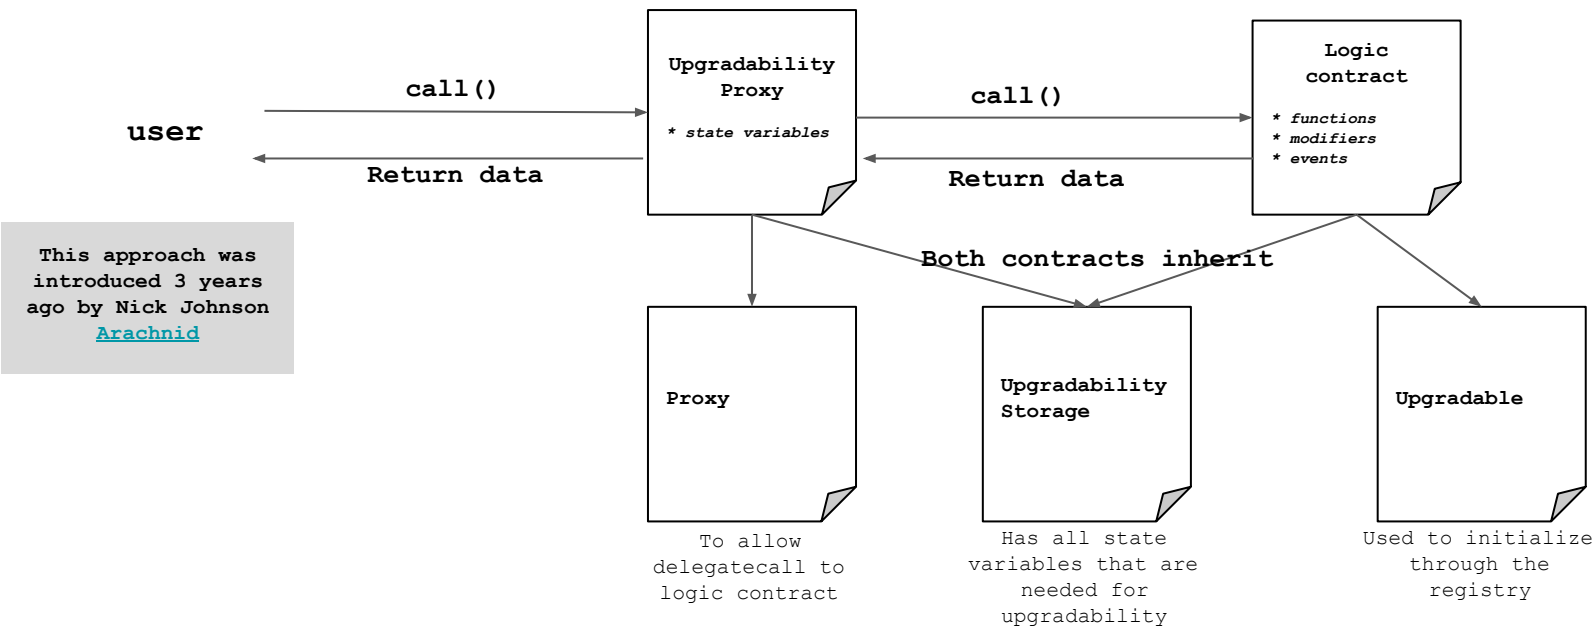
- How to perform upgrade



storage patterns		
Inherited Storage	Eternal Storage	Unstructured Storage

### Delegatecall-based proxies

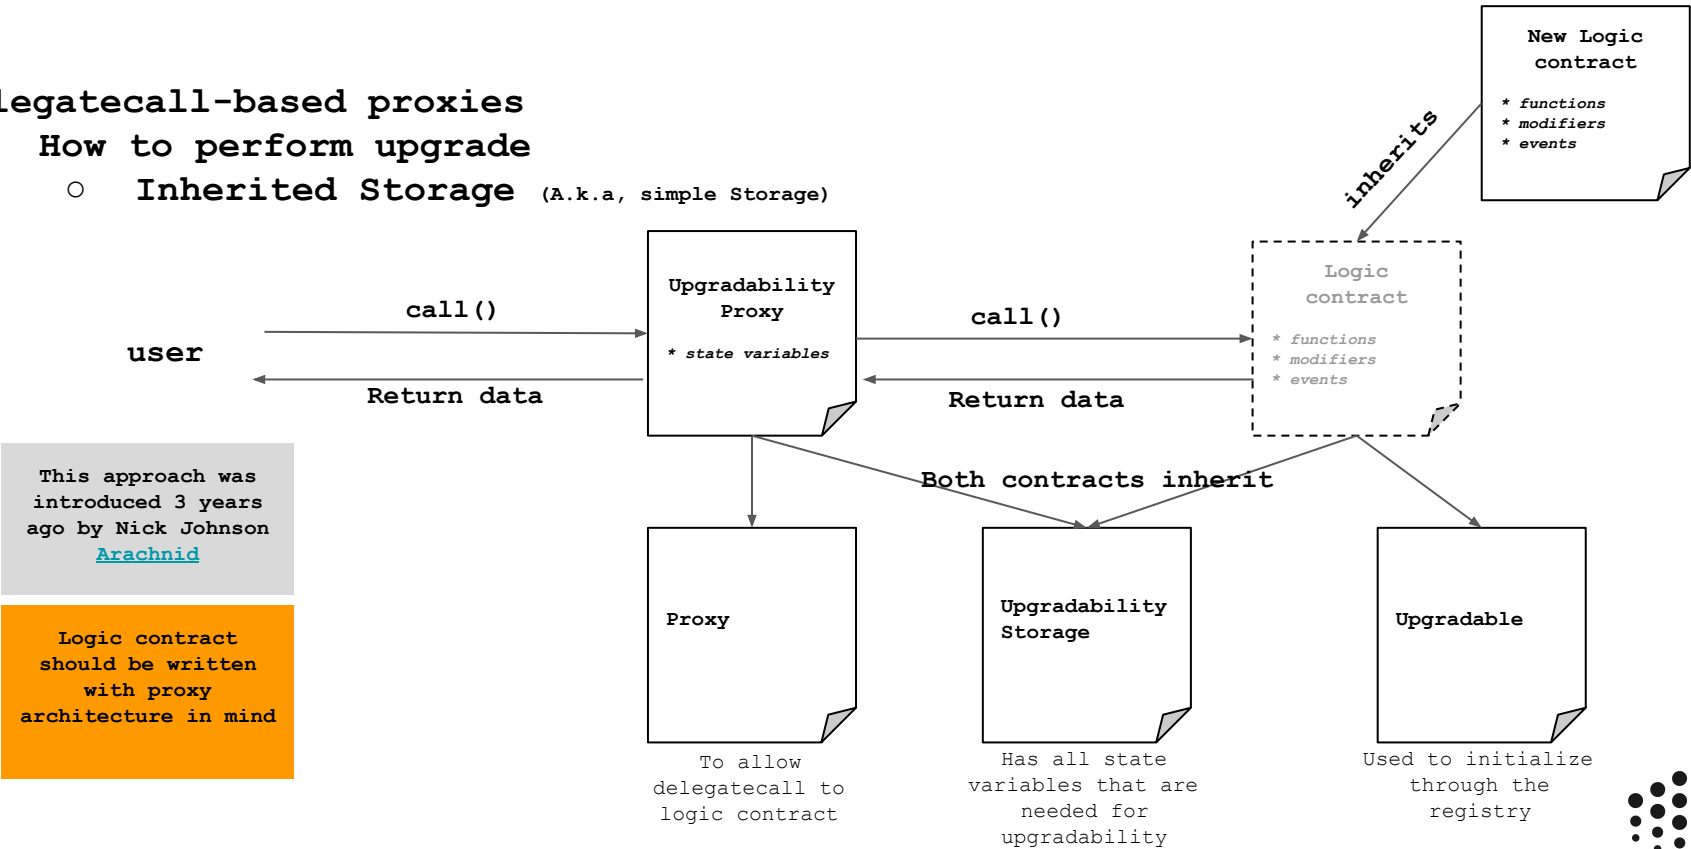
- How to perform upgrade
  - Inherited Storage (A.k.a, simple Storage)



This approach was introduced 3 years ago by Nick Johnson [Arachnid](#)

### Delegatecall-based proxies

- How to perform upgrade
  - Inherited Storage (A.k.a, simple Storage)



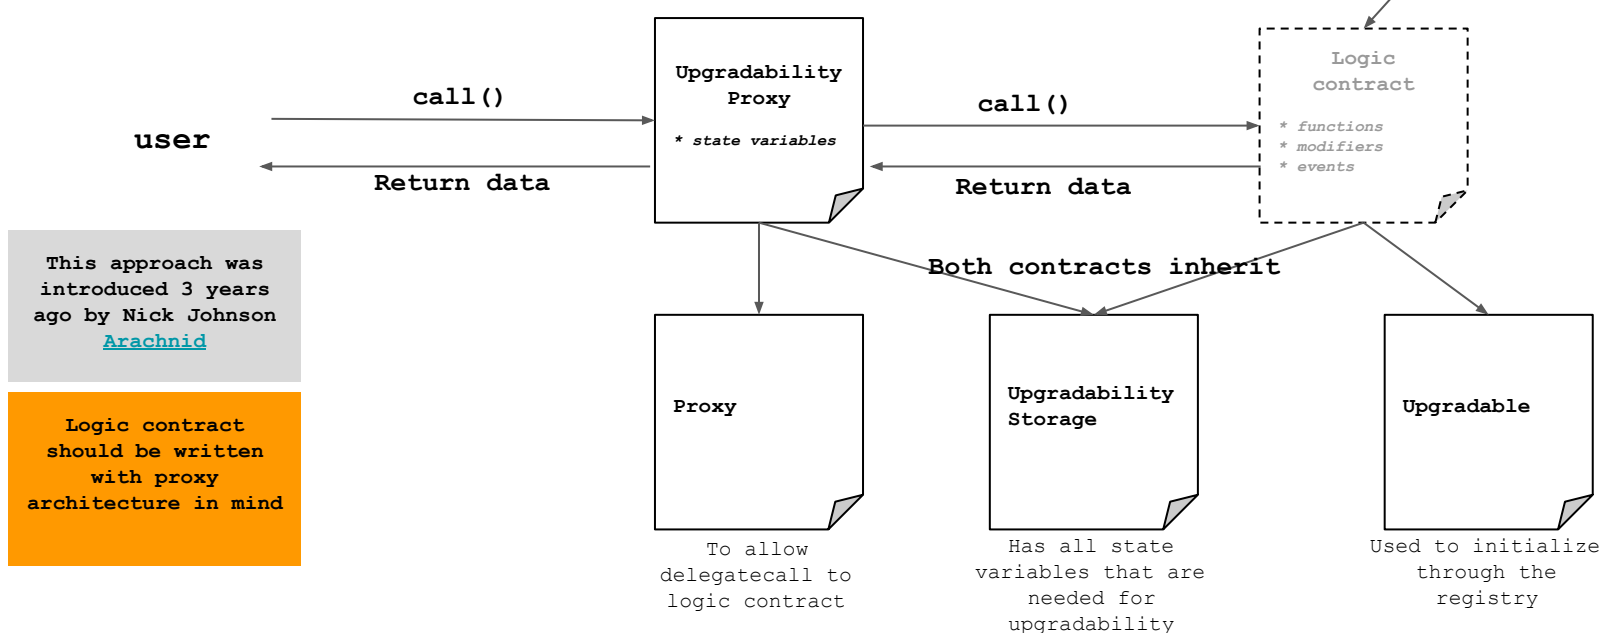
This approach was introduced 3 years ago by Nick Johnson [Arachnid](#)

Logic contract should be written with proxy architecture in mind



### Delegatecall-based proxies

- How to perform upgrade
  - Inherited Storage (A.k.a, simple Storage)

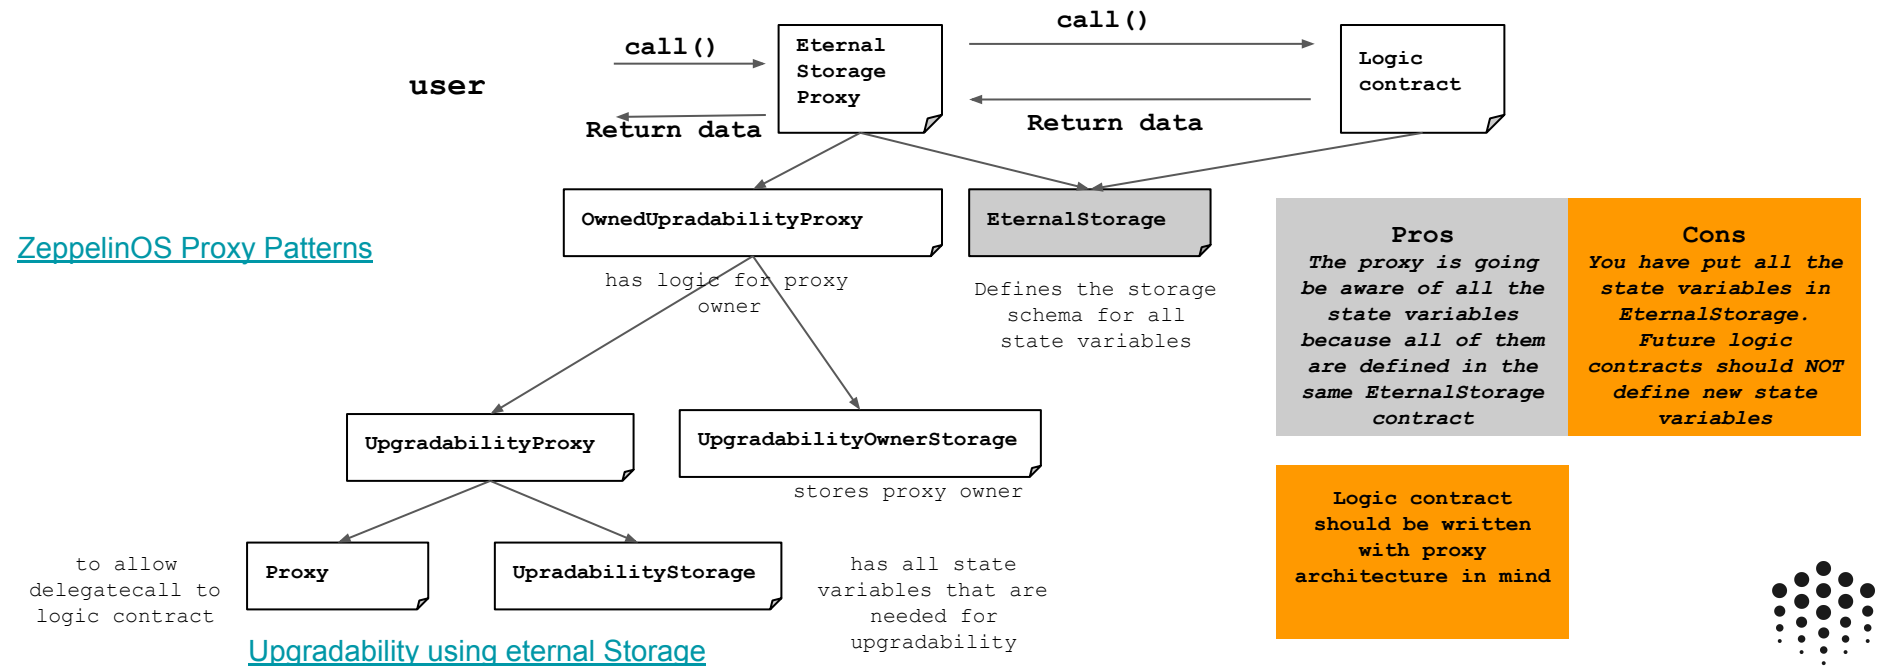


Can we have something better ?

[Upgradability using inherited storage](#)

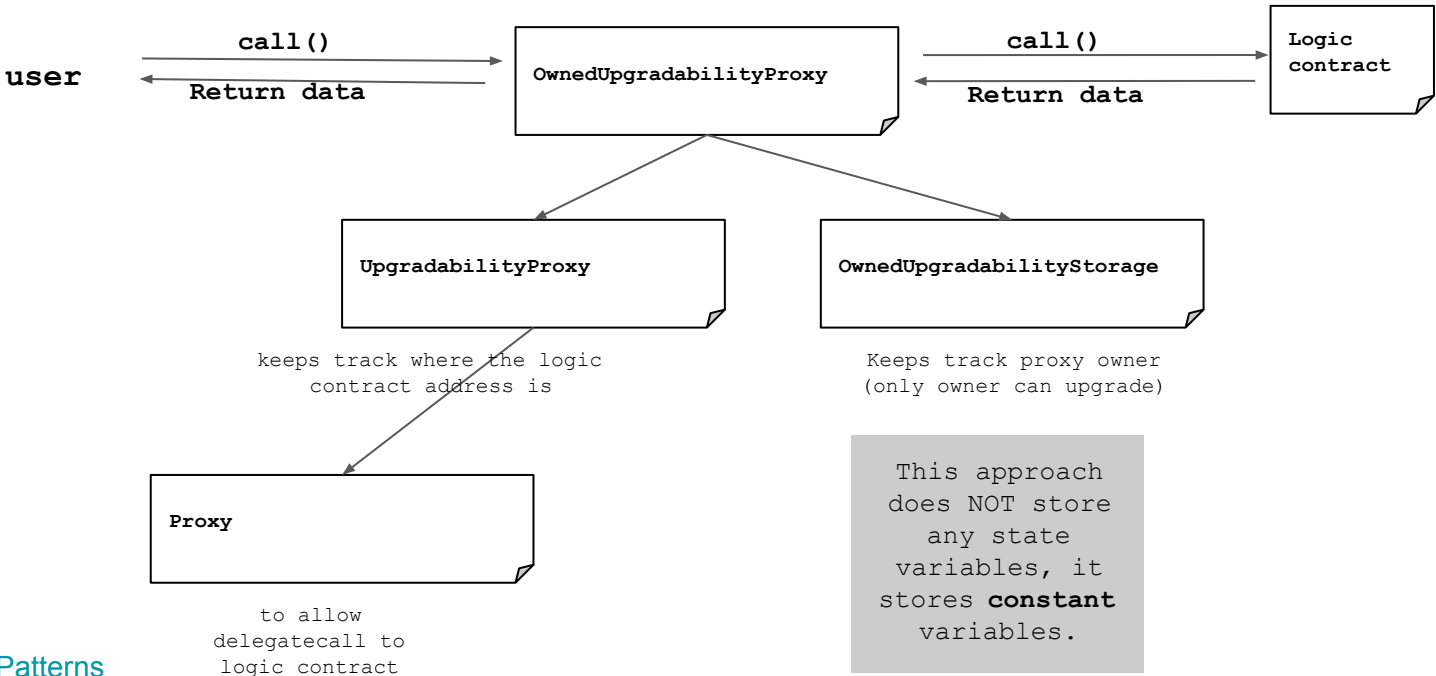
### Delegatecall-based proxies

- How to perform upgrade
  - Eternal Storage



### Delegatecall-based proxies

- How to perform upgrade
  - Unstructured Storage



### Delegatecall-based proxies

- How to perform upgrade
  - Unstructured Storage

#### ■ Logic contract:

No additional requirements!

#### Cons

*It might have a chance of hash collision for the constant position and another storage allocation. But it has a low probability*

```
contract UpgradabilityProxy is Proxy {
    bytes32 private constant implementationPosition = keccak256("com.oceanprotocol.implementation");

    function implementation()
    public
    view
    returns(address implementationAddress)
    {
        bytes32 position = implementationPosition;
        assembly {
            implementationAddress := sload(position)
        }
    }

    function setImplementation(address newImplementationAddress)
    internal
    {
        bytes32 position = implementationPosition;
        assembly{
            sstore(position, newImplementationAddress)
        }
    }
}

contract OwnedUpgradabilityProxy is UpgradabilityProxy {
    bytes32 private constant proxyOwnerPosition = keccak256("com.oceanprotocol.proxy.owner");
}
```

#### Pros

*The proxy contract does not store any state variables, therefore the logic contract can be written without any proxy architecture*

## Recommendations

- Have a detailed understanding of Ethereum internals
- Carefully consider the order of inheritance
- Carefully consider the order in which variables are declared
- Be aware that the compiler may use padding and/or pack variables together.
- Confirm that the variables' memory layout is respected
- Carefully consider the contract's initialization.
- Carefully consider names of functions in the proxy



- **Key features:**
  - New platform with support for EVM packages (on-chain standard libraries)
  - Safer and cheaper upgradeability
    - [initialization checks](#)
    - [storage layout checks](#)
    - [safe code checks](#)
  - New governance systems
    - Using [multi-signature wallet](#)

[ZeppelinOS Lab](#)



### ZeppelinOS Upgradability checklist

- ☐ You must add initializers.
- ☐ Don't forget to initialize the inherited contracts
- ☐ Make sure that all the initial values are set in an initializer function. Otherwise, any upgradable instance will not have these fields sets
- ☐ When you create new instance of an arbitrary contract inside contract, pass the instance of that contract to initialize function
- ☐ You can create new contract instances on the fly using BaseApp in ZeppelinOS
- ☐ Remember nothing prevent malicious actor from sending transactions



### ZeppelinOS Upgradability checklist

- ❑ Make sure to add the new variables at the end
- ❑ Be careful when you introduce new updates to the contracts
- ❑ Be careful with the solidity inheritance linearization
  - ❑ Avoid base contracts swapping
  - ❑ Avoid adding new variables to the base contracts.
  - ❑ A workaround for this is to declare unused variables on base contracts that you may want to extend in the future, as a means of "reserving" those slots.
  - ❑ Note that this trick does not involve increased gas usage.

Thank You!



ocean